# Test accuracy differs by at most 0.11 points
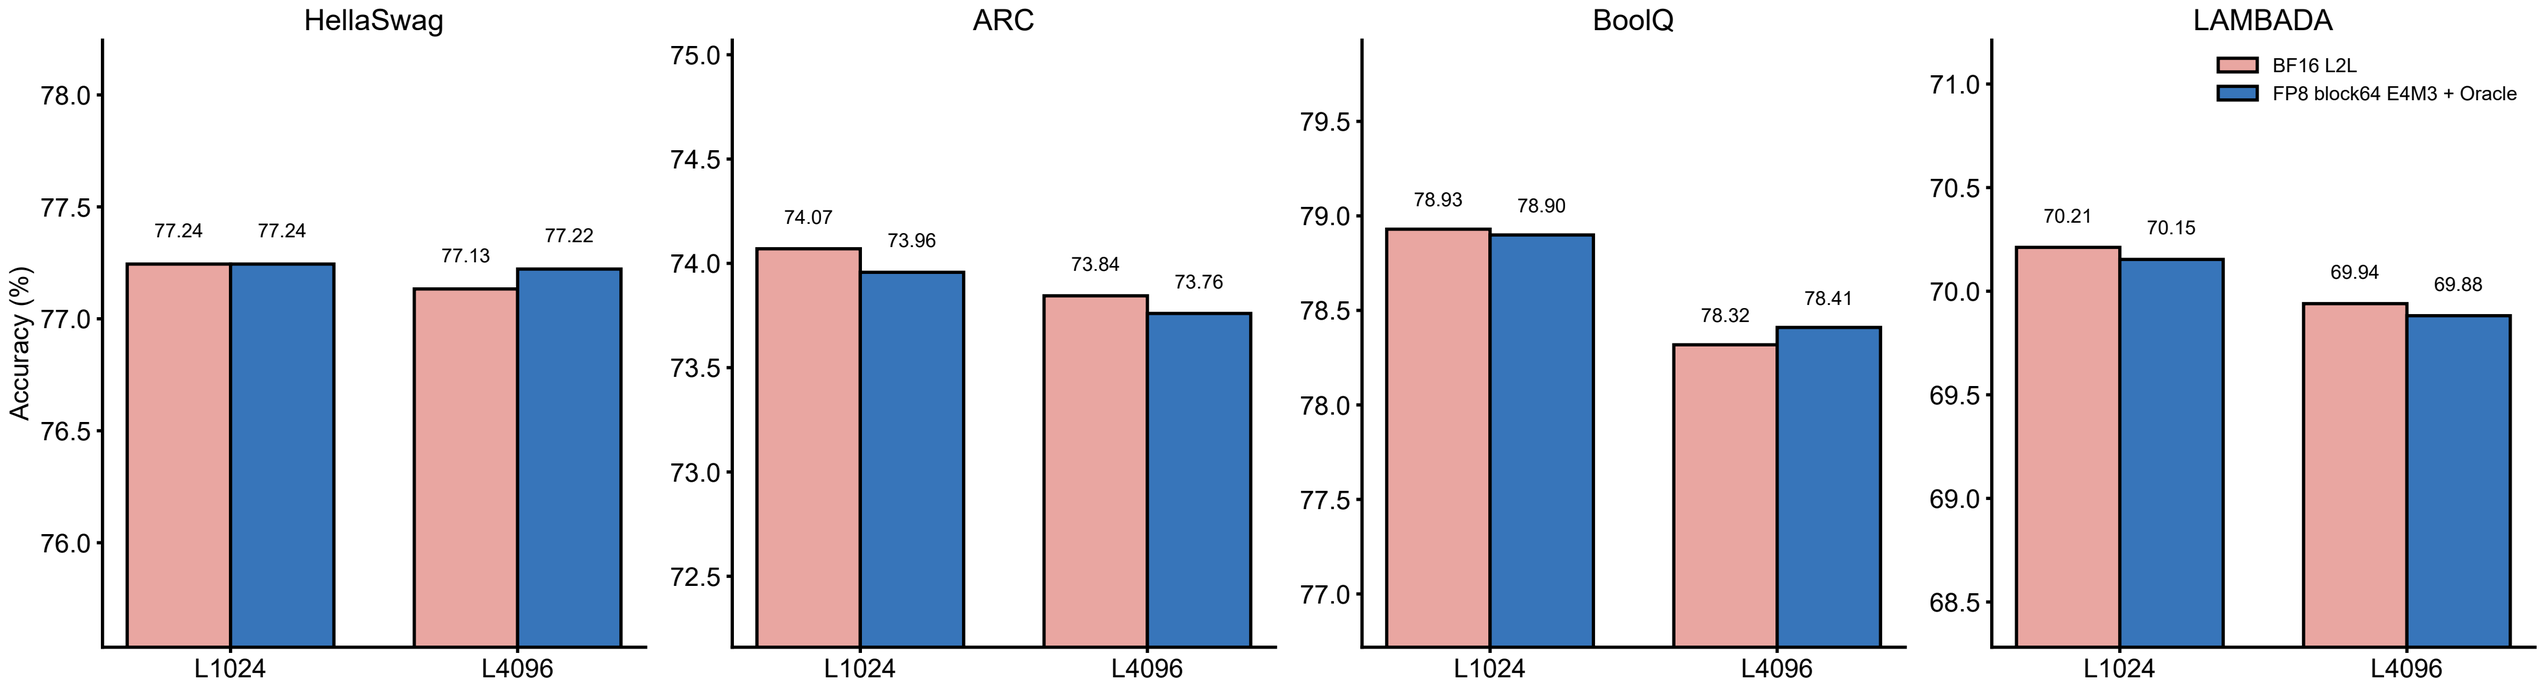
One Moonlight-16B model (mixture of experts) fine-tuned with LoRA twice for 64 steps, trained layer by layer (L2L).
Red: activations saved in BF16. Blue: saved in 8-bit FP8 (block64 E4M3), with Oracle replaying the router's expert choice in the backward pass.
Both scored with the same plain BF16 inference. Training context 1024 or 4096 tokens. Each panel's y axis is zoomed in.