Both models get the same right/wrong result on 99.1% of tasks
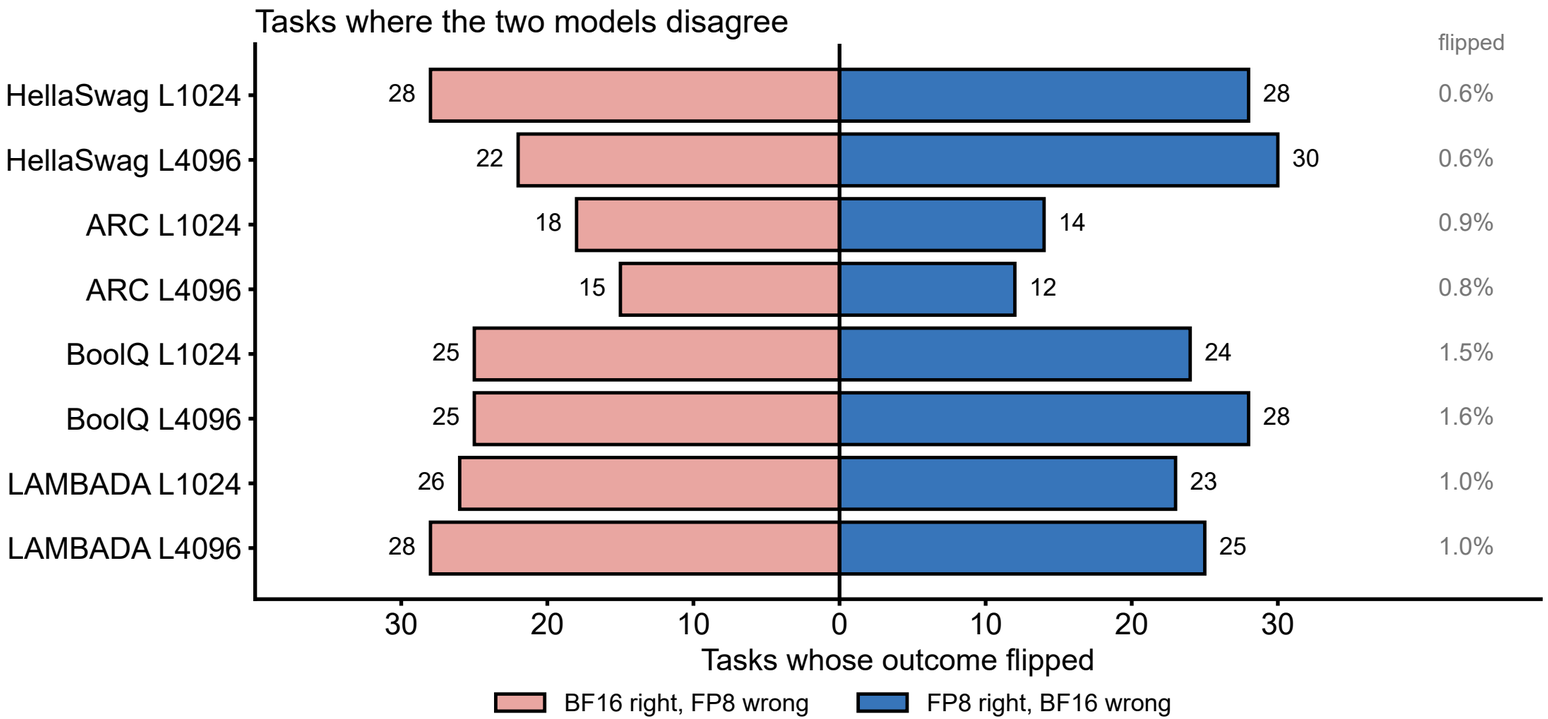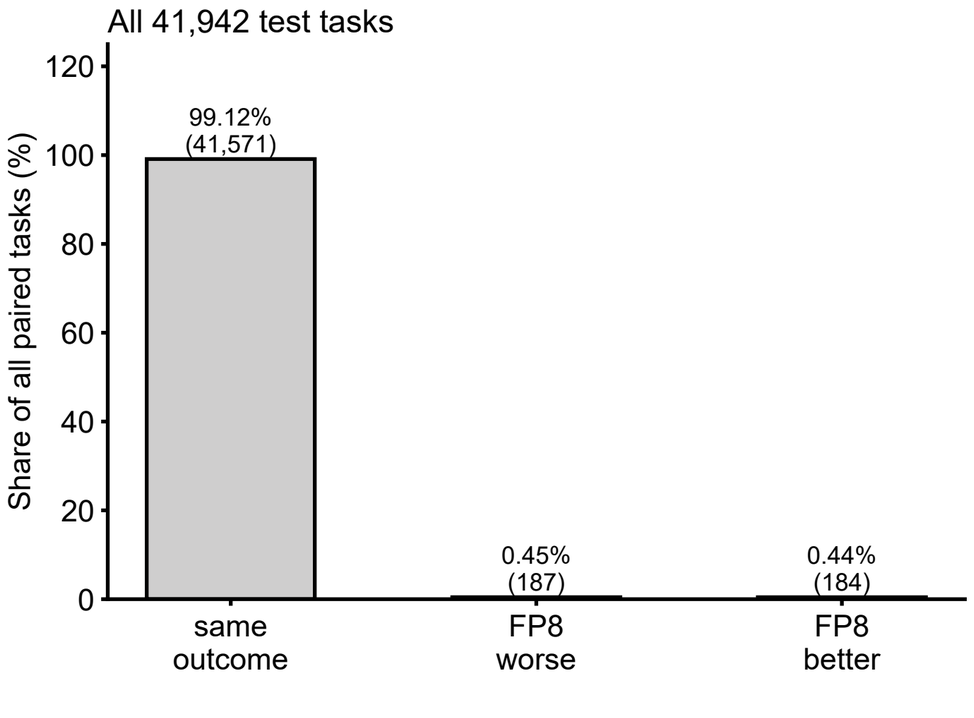
Left: per benchmark, how many tasks only BF16 solved (red) and how many only FP8 solved (blue). Gray % = share of tasks that differ.
The two sides are about equal (187 vs 184), consistent with the result in Fig. 1.