Math and code NLL moves by at most 0.22% (descriptive only)

NLL measures how surprised the model is by the reference answer (lower is better). Each dot is one problem: FP8 NLL minus BF16 NLL, per-task mean.
Black dot = mean, bars = 95% interval; bold % = change vs the BF16 mean. FP8 is slightly better on GSM8K 1024, worse on the other three.
Three of four intervals exclude zero; the shifts are small. No margin applies and these are not part of the verdict. Y axis clipped at the 99.5th percentile.

